



Module: Introduction to Parallel Programming Techniques
Module ID: EE4107
Student Name: Umid Narziev
Student ID: 200234832
Assignment Number: 7
Academic Year: 2020 - 2021

Contents

System Specification 3

TASK 1 4

 Code:..... 4

 Result: 8

 Conclusion: 8

TASK 2 9

 Code:..... 9

 Result: 11

 Conclusion: 11

TASK 3 12

 Result: 14

System Specification

Below, shown the system which was used to run and get the result of the simulation:

CPU:	Intel i5-7200U
Architecture:	Kaby Lake
Segment:	Mobile Processors
The number of cores:	2
Number of threads	4
Clock Frequency	2.50-3.10GHz (Turbo Boost)
Cache levels:	3
Cache level 1 size:	128KBytes
Cache level 2 size:	512Kbytes
Cache level 3 size:	3MBytes
RAM	12 GB
SSD:	250 GB
Operating System:	Ubuntu 20.04.2 LTS
Compiler:	Gcc and its libraries
IDE:	Clion (2020.03)

TASK 1

Code:

```
/* File:
 * 7.1.c
 *
 * Purpose:
 * Computes a parallel matrix-vector product. Matrix
 * is distributed by block rows. Vectors are distributed by
 * blocks. This version uses a random number generator to
 * generate A and x. There is some optimization.
 *
 * Compile:
 * gcc -g -Wall -fopenmp -o 7.1
 * 7.1.c
 * Run:
 * ./7.1 <thread_count> <m> <n>
 *
 * Input:
 * None unless compiled with DEBUG flag.
 * With DEBUG flag, A, x
 *
 * Output:
 * y: the product vector
 * Elapsed time for the computation
 *
 * Notes:
 * 1. Storage for A, x, y is dynamically allocated.
 * 2. Number of threads (thread_count) should evenly divide both
 * m and n. The program doesn't check for this.
 * 3. We use a 1-dimensional array for A and compute subscripts
 * using the formula  $A[i][j] = A[i*n + j]$ 
 * 4. Distribution of A, x, and y is logical: all three are
 * globally shared.
 * 5. DEBUG compile flag will prompt for input of A, x, and
 * print y
 */

#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#include "timer.h"

#define B_part
/*
 * define A_part -> for task 7.a
 * define B_part -> for task 7.b
 */

/* Serial functions */
void Get_args(int argc, char* argv[], int* thread_count_p,
              int* m_p, int* n_p);
void Usage(char* prog_name);
void Gen_matrix(double A[], int m, int n);
void Read_matrix(char* prompt, double A[], int m, int n);
void Gen_vector(double x[], int n);
void Read_vector(char* prompt, double x[], int n);
void Print_matrix(char* title, double A[], int m, int n);
void Print_vector(char* title, double y[], double m);

/* Parallel function */
void Omp_mat_vect(double A[], double x[], double y[],
```

```

        int m, int n, int thread_count);

/*-----*/
int main(int argc, char* argv[]) {
    int thread_count;
    int m, n;
    double* A;
    double* x;
    double* y;

    Get_args(argc, argv, &thread_count, &m, &n);

    A = malloc(m*n*sizeof(double));
    x = malloc(n*sizeof(double));
#ifdef A_part
    y = malloc(((8*thread_count)+m)*sizeof(double));
#endif

#ifdef B_part
    y = malloc((m)*sizeof(double));
#endif

# ifdef DEBUG
    Read_matrix("Enter the matrix", A, m, n);
    Print_matrix("We read", A, m, n);
    Read_vector("Enter the vector", x, n);
    Print_vector("We read", x, n);
# else
    Gen_matrix(A, m, n);
/* Print_matrix("We generated", A, m, n); */
    Gen_vector(x, n);
/* Print_vector("We generated", x, n); */
# endif
    double start, finish, elapsed;
    GET_TIME(start);
    Omp_mat_vect(A, x, y, m, n, thread_count);
    GET_TIME(finish);
    elapsed = finish - start;
    printf("Elapsed time = %e seconds\n", elapsed);
# ifdef DEBUG
    Print_vector("The product is", y, m);
# else
/* Print_vector("The product is", y, m); */
# endif

    free(A);
    free(x);
    free(y);

    return 0;
} /* main */

/*-----
* Function: Get_args
* Purpose:  Get command line args
* In args:  argc, argv
* Out args: thread_count_p, m_p, n_p
*/
void Get_args(int argc, char* argv[], int* thread_count_p,
              int* m_p, int* n_p) {

```

```

    if (argc != 4) Usage(argv[0]);
    *thread_count_p = strtol(argv[1], NULL, 10);
    *m_p = strtol(argv[2], NULL, 10);
    *n_p = strtol(argv[3], NULL, 10);
    if (*thread_count_p <= 0 || *m_p <= 0 || *n_p <= 0) Usage(argv[0]);

} /* Get_args */

/*-----
 * Function: Usage
 * Purpose:  print a message showing what the command line should
 *           be, and terminate
 * In arg :  prog_name
 */
void Usage (char* prog_name) {
    fprintf(stderr, "usage: %s <thread_count> <m> <n>\n", prog_name);
    exit(0);
} /* Usage */

/*-----
 * Function:  Read_matrix
 * Purpose:   Read in the matrix
 * In args:   prompt, m, n
 * Out arg:   A
 */
void Read_matrix(char* prompt, double A[], int m, int n) {
    int i, j;

    printf("%s\n", prompt);
    for (i = 0; i < m; i++)
        for (j = 0; j < n; j++)
            scanf("%lf", &A[i*n+j]);
} /* Read_matrix */

/*-----
 * Function: Gen_matrix
 * Purpose:  Use the random number generator random to generate
 *           the entries in A
 * In args:  m, n
 * Out arg:  A
 */
void Gen_matrix(double A[], int m, int n) {
    int i, j;
    for (i = 0; i < m; i++)
        for (j = 0; j < n; j++)
            A[i*n+j] = random()/((double) RAND_MAX);
} /* Gen_matrix */

/*-----
 * Function: Gen_vector
 * Purpose:  Use the random number generator random to generate
 *           the entries in x
 * In arg:   n
 * Out arg:  A
 */
void Gen_vector(double x[], int n) {
    int i;
    for (i = 0; i < n; i++)
        x[i] = random()/((double) RAND_MAX);
} /* Gen_vector */

/*-----

```

```

* Function:    Read_vector
* Purpose:     Read in the vector x
* In arg:      prompt, n
* Out arg:     x
*/
void Read_vector(char* prompt, double x[], int n) {
    int i;

    printf("%s\n", prompt);
    for (i = 0; i < n; i++)
        scanf("%lf", &x[i]);
} /* Read_vector */

/*-----
* Function: Omp_mat_vect
* Purpose:  Multiply an mxn matrix by an nx1 column vector
* In args:  A, x, m, n, thread_count
* Out arg:  y
*/
void Omp_mat_vect(double A[], double x[], double y[],
                  int m, int n, int thread_count) {
    int i, j;
    double temp;

    # pragma omp parallel for num_threads(thread_count) default(none) private(i, j, temp) shared(A, x, y, m, n)
    for (i = 0; i < m; i++) {
        int my_rank = omp_get_thread_num();
    #ifdef A_part
        y[i+(my_rank*8)] = 0.0;
    #endif
    #ifdef B_Part
        y[i] = 0.0;
        temp = 0.0;
    #endif
        for (j = 0; j < n; j++) {
    #ifdef A_part
        y[i+(my_rank*8)] += A[i*n+j]*x[j];;
    #endif
    #ifdef B_part
        temp += A[i*n+j]*x[j];
    #endif
        }
    #ifdef B_part
        y[i] = temp;
    #endif
    }

} /* Omp_mat_vect */

/*-----
* Function: Print_matrix
* Purpose:  Print the matrix
* In args:  title, A, m, n
*/
void Print_matrix(char* title, double A[], int m, int n) {
    int i, j;

```

```

printf("%s\n", title);
for (i = 0; i < m; i++) {
    for (j = 0; j < n; j++)
        printf("%4.1f ", A[i*n + j]);
    printf("\n");
}
} /* Print_matrix */

/*-----
 * Function:  Print_vector
 * Purpose:   Print a vector
 * In args:   title, y, m
 */
void Print_vector(char* title, double y[], double m) {
    int i;

    printf("%s\n", title);
    for (i = 0; i < m; i++)
        printf("%4.1f ", y[i]);
    printf("\n");
} /* Print_vector */

```

Result:

A) Padded

```

umid@umid-Lenovo-Ideapad-320-15IKB:/media/umid/Data/Aston_University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.1$ ./7.a 4 8 8000000
Elapsed time = 1.285881e-01 seconds

```

B) With private storage

```

umid@umid-Lenovo-Ideapad-320-15IKB:/media/umid/Data/Aston_University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.1$ ./7.a 4 8 8000000
Elapsed time = 9.817314e-02 seconds

```

C) Original program

```

umid@umid-Lenovo-Ideapad-320-15IKB:/media/umid/Data/Aston_University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.1$ ./7.1 4 8 8000000
Elapsed time = 1.985300e-01 seconds

```

Conclusion:

- In fact, the original program has the worst performance among the three.
- The padded version of the program did not show the best result, however, solid improvement is seen, whereas the private variable program has shown a huge improvement.

TASK 2

Code:

```
/*
 * To compile -> gcc -g -Wall -fopenmp -o 7.2 7.2.c
 * To run -> ./7.2 thread_number number_of_element
 */
#include <stdio.h>
#include <string.h>
#include <malloc.h>
#include <stdlib.h>
#include <time.h>
#include "omp.h"

// #define PRINT // To see the elements

void count_Sort_Serial(int a[], int n);
void count_Sort_Parallel(int a[], int n);
void print(int a[], int n);
void usage(char* prog_name);
void array_Generating(int a[], int n);
int comparator(const void * p1, const void * p2);
int thread_count;

int main(int argc, char* argv[]) {

    int n; /// number of elements
    int *a, *b, *c; /// pointer to the array
    double start, end, temp_Parallel, temp_Serial, temp_qSort; /// timing variables

    if (argc != 3) usage(argv[0]);
    thread_count = strtol(argv[1], NULL, 10);
    n = strtol(argv[2], NULL, 10);
    a = malloc(n * sizeof(int));
    b = malloc(n * sizeof(int));
    c = malloc(n * sizeof(int));

    array_Generating(a, n);
    memcpy(b, a, n * sizeof(int)); // Duplicating the values
    memcpy(c, a, n * sizeof(int)); // Duplicating the values
#ifdef PRINT
    print(a, n);
#endif
    //-----Parallel Count Sort-----
    start = omp_get_wtime();
    count_Sort_Parallel(a, n);
    end = omp_get_wtime();
#ifdef PRINT
    printf("After Parallel Count Sort function: \n");
    print(a, n);
#endif
    temp_Parallel = end - start;
    //-----Serial Count Sort-----
    start = omp_get_wtime();
    count_Sort_Serial(b, n);
    end = omp_get_wtime();
#ifdef PRINT
    printf("After Serial Count Sort function: \n");
    print(b, n);
#endif
    temp_Serial = end - start;
```

```

//-----qSort Count Sort-----
start = omp_get_wtime();
qsort(c, n, sizeof(int), comparator);
end = omp_get_wtime();
#ifdef PRINT
    printf("After qSort Count Sort function: \n");
    print(b, n);
#endif
temp_qSort = end - start;
//-----The result printing -----
printf("The elapsed time is %e for Parallel count sort function with %d elements.\n", temp_Parallel, n);
printf("The elapsed time is %e for Serial count sort function with %d elements.\n", temp_Serial, n);
printf("The elapsed time is %e for Qsort count sort function with %d elements.\n", temp_qSort, n);
return 0;
}

void usage(char* prog_name) {

    fprintf(stderr, "usage: %s <number of threads>, <number of elements>\n", prog_name);
    exit(0);
}

void array_Generating(int a[], int n){

    for (int i = 0; i < n; ++i) {
        srand(i);
        a[i] = rand() % n;
    }
}

void print(int a[], int n){
    for (int i = 0; i < n; i++) {
        printf("%d\t", a[i]);
    }
    printf("\n");
}

void count_Sort_Serial(int a[], int n){
    int count;
    int *temp = malloc(n*sizeof(int));

    for (int i = 0; i < n; ++i) {
        count = 0;
        for (int j = 0; j < n; ++j) {
            if(a[j]<a[i]){
                count++;
            } else if((a[j] == a[i]) && (j < i)){
                count++;
            }
        }
        temp[count] = a[i];
    }

    memcpy(a, temp, n * sizeof(int));
    free(temp);
}

void count_Sort_Parallel(int a[], int n){
    int count, i, j;
    int *temp = malloc(n*sizeof(int));
#pragma omp parallel for num_threads(thread_count) default(none) shared(a, n, temp) private(j, i, count)
    for (i = 0; i < n; ++i) {
        count = 0;
        for (j = 0; j < n; ++j) {
            if(a[j] < a[i]){
                count++;
            }
        }
        temp[count] = a[i];
    }
}

```

```

        } else if((a[j] == a[i]) && (j < i)){
            count++;
        }
    }
#pragma omp critical
    {
        temp[count] = a[i];
    };

}

memcpy(a, temp, n * sizeof(int));

free(temp);
}
int comparator (const void * p1, const void * p2){
    return (*(int*)p1 - *(int*)p2);
}

```

Result:

```

umid@umid-Lenovo-ideapad-320-15IKB:/media/umid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.2$ gcc -g -Wall -fopenmp -o 7.2 7.2.c
umid@umid-Lenovo-ideapad-320-15IKB:/media/umid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.2$ ./7.2 4 1000
The elapsed time is 1.217650e-02 for Parallel count sort function with 1000 elements.
umid@umid-Lenovo-ideapad-320-15IKB:/media/umid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.2$ ./7.2 4 1000
The elapsed time is 1.302248e-02 for Serial count sort function with 1000 elements.
umid@umid-Lenovo-ideapad-320-15IKB:/media/umid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.2$ ./7.2 4 1000
The elapsed time is 1.743030e-04 for Qsort count sort function with 1000 elements.

```

Conclusion:

- The private variables are array j, i, count, shared variables are array a, n, temp.
- No, there is no loop cared dependency as i-private.
- The memcpy function is not IO-bound, not CPU-bound, so parallelizing will not give much performance boost. And also, parallelizing it not safe as we don't know how it is implemented, memory location may be overwritten.
- The qSort function has the highest performance among all, the parallelising of count sort has given benefit but not as much as qsort.

TASK 3

Code:

```
/*
 * To compile -> gcc -g -Wall -fopenmp -o 7.3 7.3.c
 * To run -> ./7.3
 * export OMP_SCHEDULE="guided, 500"
 */

#include <stdio.h>
#include <malloc.h>
#include <stdlib.h>
#include "omp.h"

void row_Oriented_Inner(double A[], double b[], double x[], int n, int thread_count);
void column_Oriented_Inner(double A[], double b[], double x[], int n, int thread_count);
void generatorMatrixVector(double A[], double b[], int n);
void matrixVectorPrint(double A[], double b[], int n);
void printVector(double x[], int n);
int main() {
    int n = 10000;
    double *a = malloc(n * n * sizeof(double));
    double *b = malloc(n * n * sizeof(double));
    double *x = malloc(n * sizeof(double));
    double *y = malloc(n * sizeof(double));
    double start, end;
    generatorMatrixVector(a, b, n);
    for (int thread_count = 1; thread_count < 16; thread_count<=1) {
        start = omp_get_wtime();
        column_Oriented_Inner(a, b, x, n, thread_count);
        end = omp_get_wtime();
        printf("[Column] The elapsed time -> %f, thread_count -> %d, n -> %d, thread_count->%d\n", end - start,
            thread_count, n, thread_count);
    }
    for (int thread_count = 1; thread_count < 16; thread_count<=1) {
        start = omp_get_wtime();
        row_Oriented_Inner(a, b, y, n, thread_count);
        end = omp_get_wtime();
        printf("[Row] The elapsed time -> %f, thread_count -> %d, n -> %d, thread_count->%d\n", end - start,
            thread_count, n, thread_count);
    }
    return 0;
}

void row_Oriented_Inner(double A[], double b[], double x[], int n, int thread_count){
    double temp;
    int row, col;
    #pragma omp parallel default(none) num_threads(thread_count) \
        shared(n, b, A, temp, x) private(col, row)
        for ( row = n-1; row >= 0; row--) {
    #pragma omp single
        temp = b[row];
    #pragma omp for reduction(+:temp) schedule(runtime)
        for ( col = row + 1; col < n; col++) {
            temp += -A[row*n+col] * x[col];
        }
    #pragma omp single // using critical gives strange result as it is distributed over process so I used again
    single
```

```

        x[row] = temp / A[row*n+row] ;
    }
}

void column_Oriented_Inner(double A[], double b[], double x[], int n, int thread_count){
    int col, row;
    #pragma omp parallel default(none) num_threads(thread_count) \
        shared(n, b, A, x) private(col, row)
    {
        #pragma omp for
        for (int row = 0; row < n; ++row) {
            x[row] = b[row];
        }
        for (col = n - 1; col >= 0; col--) {
            #pragma omp single
            x[col] /= A[col*n+col];
            #pragma omp for schedule(runtime)
            for (row = 0; row < col; row++) {
                x[row] +=- A[row*n+col] * x[col];
            }
        }
    }
}

void generatorMatrixVector(double A[], double b[], int n){
    srand(1);
    for (int col = 0; col < n; ++col) {
        for (int row = 0; row < n; ++row) {
            A[row*n+col] = random()/((double) RAND_MAX);
        }
    }
    srand(2);
    for (int i = 0; i < n; ++i) {
        b[i] = random()/((double) RAND_MAX);
    }
}

void matrixVectorPrint(double A[], double b[], int n){

    printf("The generated matrix\n");
    for (int col = 0; col < n; ++col) {
        for (int row = 0; row < n; ++row) {
            printf("%f\t", A[row*n+col]);
        }
        printf("\n");
    }
    printf("The generated vector\n");
    for (int i = 0; i < n; ++i) {
        printf("%f\t", b[i]);
    }
    printf("\n");
}

void printVector(double x[], int n){

    printf("The vector is: \n");
    for (int i = 0; i < n; ++i) {
        printf("%10.4f\t", x[i]);
    }
    printf("\n");
}

```

Result:

1. Default Scheduler:

```
uid@uid-Lenovo-ideapad-320-15IKB:/media/uid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.3$ gcc -g -Wall -fopenmp -o 7.3 7.3.c
uid@uid-Lenovo-ideapad-320-15IKB:/media/uid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.3$ ./7.3
[Column] The elapsed time -> 2.805761, thread_count -> 1, n -> 10000, thread_count->1
[Column] The elapsed time -> 2.751768, thread_count -> 2, n -> 10000, thread_count->2
[Column] The elapsed time -> 1.678471, thread_count -> 4, n -> 10000, thread_count->4
[Column] The elapsed time -> 1.801062, thread_count -> 8, n -> 10000, thread_count->8
[Row] The elapsed time -> 0.900965, thread_count -> 1, n -> 10000, thread_count->1
[Row] The elapsed time -> 1.440938, thread_count -> 2, n -> 10000, thread_count->2
[Row] The elapsed time -> 0.982260, thread_count -> 4, n -> 10000, thread_count->4
[Row] The elapsed time -> 1.605797, thread_count -> 8, n -> 10000, thread_count->8
```

2. Static, chunk size = 8:

```
uid@uid-Lenovo-ideapad-320-15IKB:/media/uid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.3$ export OMP_SCHEDULE="static, 8"
uid@uid-Lenovo-ideapad-320-15IKB:/media/uid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.3$ ./7.3
[Column] The elapsed time -> 0.684080, thread_count -> 1, n -> 10000, thread_count->1
[Column] The elapsed time -> 0.435846, thread_count -> 2, n -> 10000, thread_count->2
[Column] The elapsed time -> 0.384885, thread_count -> 4, n -> 10000, thread_count->4
[Column] The elapsed time -> 0.813419, thread_count -> 8, n -> 10000, thread_count->8
[Row] The elapsed time -> 0.177707, thread_count -> 1, n -> 10000, thread_count->1
[Row] The elapsed time -> 0.172560, thread_count -> 2, n -> 10000, thread_count->2
[Row] The elapsed time -> 0.141851, thread_count -> 4, n -> 10000, thread_count->4
[Row] The elapsed time -> 0.641442, thread_count -> 8, n -> 10000, thread_count->8
```

3. Guided, chunk size = 8:

```
uid@uid-Lenovo-ideapad-320-15IKB:/media/uid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.3$ export OMP_SCHEDULE="guided, 8"
uid@uid-Lenovo-ideapad-320-15IKB:/media/uid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.3$ ./7.3
[Column] The elapsed time -> 0.672911, thread_count -> 1, n -> 10000, thread_count->1
[Column] The elapsed time -> 0.394495, thread_count -> 2, n -> 10000, thread_count->2
[Column] The elapsed time -> 0.347416, thread_count -> 4, n -> 10000, thread_count->4
[Column] The elapsed time -> 0.632058, thread_count -> 8, n -> 10000, thread_count->8
[Row] The elapsed time -> 0.178246, thread_count -> 1, n -> 10000, thread_count->1
[Row] The elapsed time -> 0.111092, thread_count -> 2, n -> 10000, thread_count->2
[Row] The elapsed time -> 0.162991, thread_count -> 4, n -> 10000, thread_count->4
[Row] The elapsed time -> 0.497792, thread_count -> 8, n -> 10000, thread_count->8
```

4. Dynamic, chunk size = 8:

```
uid@uid-Lenovo-ideapad-320-15IKB:/media/uid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.3$ export OMP_SCHEDULE="dynamic, 8"
uid@uid-Lenovo-ideapad-320-15IKB:/media/uid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.3$ ./7.3
[Column] The elapsed time -> 0.927137, thread_count -> 1, n -> 10000, thread_count->1
[Column] The elapsed time -> 0.577075, thread_count -> 2, n -> 10000, thread_count->2
[Column] The elapsed time -> 0.576333, thread_count -> 4, n -> 10000, thread_count->4
[Column] The elapsed time -> 0.681261, thread_count -> 8, n -> 10000, thread_count->8
[Row] The elapsed time -> 0.278211, thread_count -> 1, n -> 10000, thread_count->1
[Row] The elapsed time -> 0.230488, thread_count -> 2, n -> 10000, thread_count->2
[Row] The elapsed time -> 0.190849, thread_count -> 4, n -> 10000, thread_count->4
[Row] The elapsed time -> 0.575282, thread_count -> 8, n -> 10000, thread_count->8
```

5. Auto:

```
uid@uid-Lenovo-ideapad-320-15IKB:/media/uid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.3$ export OMP_SCHEDULE="auto"
uid@uid-Lenovo-ideapad-320-15IKB:/media/uid/Data/Aston University/Subjects/TP2/EE4107 - Introduction to Parallel Programming Techniques/Assignments/Assignment - 7/7.3$ ./7.3
[Column] The elapsed time -> 0.679225, thread_count -> 1, n -> 10000, thread_count->1
[Column] The elapsed time -> 0.350997, thread_count -> 2, n -> 10000, thread_count->2
[Column] The elapsed time -> 0.330308, thread_count -> 4, n -> 10000, thread_count->4
[Column] The elapsed time -> 0.782948, thread_count -> 8, n -> 10000, thread_count->8
[Row] The elapsed time -> 0.179200, thread_count -> 1, n -> 10000, thread_count->1
[Row] The elapsed time -> 0.102238, thread_count -> 2, n -> 10000, thread_count->2
[Row] The elapsed time -> 0.101052, thread_count -> 4, n -> 10000, thread_count->4
[Row] The elapsed time -> 0.668183, thread_count -> 8, n -> 10000, thread_count->8
```

Conclusion:

- No, outer loop of the row-oriented algorithm can not be parallelized, $x[\text{row}] = b[\text{row}]$ and $x[\text{row}]/=A[\text{row}][\text{row}]$, introduce loop carried dependency.
- Yes, the inner loop of the row-oriented algorithm can be parallelized.
- No, outer loop of column-oriented algorithm can not be parallelized as $x[\text{col}]/=A[\text{col}][\text{col}]$ introduce loop-dependency.
- Yes, the inner loop of the column-oriented algorithm can be parallelized, and code is written.
- The guided scheduler has performance, with ~0.63 seconds for row-oriented algorithm and ~0.49 seconds for the column-oriented algorithm for 8 threads.